

# Dataset\_Stoke

May 30, 2021

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib
import matplotlib.ticker as mtick
import seaborn as sns
sns.set_style('white')
import plotly.express as px
import plotly.graph_objs as pgo
import plotly.offline as pyo
from plotly.subplots import make_subplots
import plotly.figure_factory as ff
pyo.init_notebook_mode()
from imblearn.over_sampling import SMOTE
import scikitplot as skplt

import joblib

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score

from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
import os

from sklearn.metrics import
    ↪ classification_report, confusion_matrix, accuracy_score, recall_score,
    ↪ precision_score, f1_score
import warnings
warnings.filterwarnings('ignore')

plt.rc('figure', figsize=(17, 13))
```

```
sns.set_context('paper',font_scale=2)

def set_seed(seed=20210524):
    np.random.seed(seed)
    os.environ['PYTHONHASHSEED'] = str(seed)
    os.environ['TF_DETERMINISTIC_OPS'] = '1'
set_seed()
```

# 1 Step one: Handle null data

## 1.1 have a glance at the data

```
[2]: df_stroke = pd.read_csv("./data/healthcare-dataset-stroke-data.csv")
df_stroke.head()
```

```
[2]:
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Self-employed	Rural	202.21	NaN	never smoked	
2	Private	Rural	105.92	32.5	never smoked	
3	Private	Urban	171.23	34.4	smokes	
4	Self-employed	Rural	174.12	24.0	never smoked	

	stroke
0	1
1	1
2	1
3	1
4	1

```
[3]: df_stroke.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    5110 non-null  int64
1   gender                5110 non-null  object
2   age                   5110 non-null  float64
3   hypertension           5110 non-null  int64
4   heart_disease          5110 non-null  int64
```

```

5  ever_married      5110 non-null  object
6  work_type         5110 non-null  object
7  Residence_type    5110 non-null  object
8  avg_glucose_level 5110 non-null  float64
9  bmi               4909 non-null  float64
10 smoking_status    5110 non-null  object
11 stroke            5110 non-null  int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB

```

## 1.2 Find null datas

```
[4]: df_stroke.isnull().sum()
```

```

[4]: id                0
     gender            0
     age              0
     hypertension      0
     heart_disease     0
     ever_married      0
     work_type         0
     Residence_type    0
     avg_glucose_level 0
     bmi              201
     smoking_status    0
     stroke            0
     dtype: int64

```

## 1.3 To deal with NaN data, I want to evaluate these bmi data with age and gender, using decision tree

```

[5]: bmi_pipe = Pipeline( steps = [
      ('scaling',StandardScaler()),
      ('lr',DecisionTreeRegressor(random_state = 30))
    ] )

cp = df_stroke[['age','gender','bmi']].copy()
cp.gender = cp.gender.replace({'Male':0,'Female':1,'Other':-1}).astype(np.uint8)

miss_bmi = cp[cp.bmi.isnull()]
cp = cp[~cp.bmi.isnull()]
bmi = cp.pop('bmi')
bmi_pipe.fit(cp,bmi)
predict_bmi = pd.Series(bmi_pipe.predict(miss_bmi[['age','gender']] ),index =
↳miss_bmi.index)
df_stroke.loc[miss_bmi.index,'bmi'] = predict_bmi

```

having a look at the predicted bmi data

```
[6]: predict_bmi
```

```
[6]: 1      29.879487
      8      30.556098
      13     27.247222
      19     30.841860
      27     33.146667
      ...
     5039    32.716000
     5048    28.313636
     5093    31.459322
     5099    28.313636
     5105    28.476923
      Length: 201, dtype: float64
```

## 1.4 so now we have all the datas

```
[7]: df_stroke.isnull().sum()
```

```
[7]: id                0
     gender            0
     age              0
     hypertension      0
     heart_disease      0
     ever_married       0
     work_type         0
     Residence_type     0
     avg_glucose_level  0
     bmi               0
     smoking_status     0
     stroke            0
     dtype: int64
```

## 2 Step two: Now, let's find the relation ship among every predictors

### 2.1 change these string value to distriected integers

```
[8]: df_stroke_int = df_stroke.copy()
     df_stroke_int.gender = df_stroke_int.gender.replace({'Male':0, 'Female':
     ↪1, 'Other':-1}).astype(np.uint8)
     df_stroke_int.smoking_status = df_stroke_int.smoking_status.replace({'formerly_
     ↪smoked':5, 'never smoked':0, 'smokes':3, 'Unknown':-1}).astype(np.uint8)
     df_stroke_int.work_type = df_stroke_int.work_type.replace({'children':
     ↪4, 'Govt_job':3, 'Never_worked':2, 'Self-employed':1, 'Private':0}).astype(np.
     ↪uint8)
```

```
df_stroke_int.ever_married = df_stroke_int.ever_married.replace({'Yes':1, 'No':
↪0}).astype(np.uint8)
df_stroke_int.Residence_type = df_stroke_int.Residence_type.replace({'Rural':
↪1, 'Urban':0}).astype(np.uint8)
df_stroke_int
```

```
[8]:
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	0	67.0	0	1	1	
1	51676	1	61.0	0	0	1	
2	31112	0	80.0	0	1	1	
3	60182	1	49.0	0	0	1	
4	1665	1	79.0	1	0	1	
...	...	...	...	...	...	...	
5105	18234	1	80.0	1	0	1	
5106	44873	1	81.0	0	0	1	
5107	19723	1	35.0	0	0	1	
5108	37544	0	51.0	0	0	1	
5109	44679	1	44.0	0	0	1	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	0	0	228.69	36.600000	5	
1	1	1	202.21	29.879487	0	
2	0	1	105.92	32.500000	0	
3	0	0	171.23	34.400000	3	
4	1	1	174.12	24.000000	0	
...	...	...	...	...	...	
5105	0	0	83.75	28.476923	0	
5106	1	0	125.20	40.000000	0	
5107	1	1	82.99	30.600000	0	
5108	0	1	166.29	25.600000	5	
5109	3	0	85.28	26.200000	255	

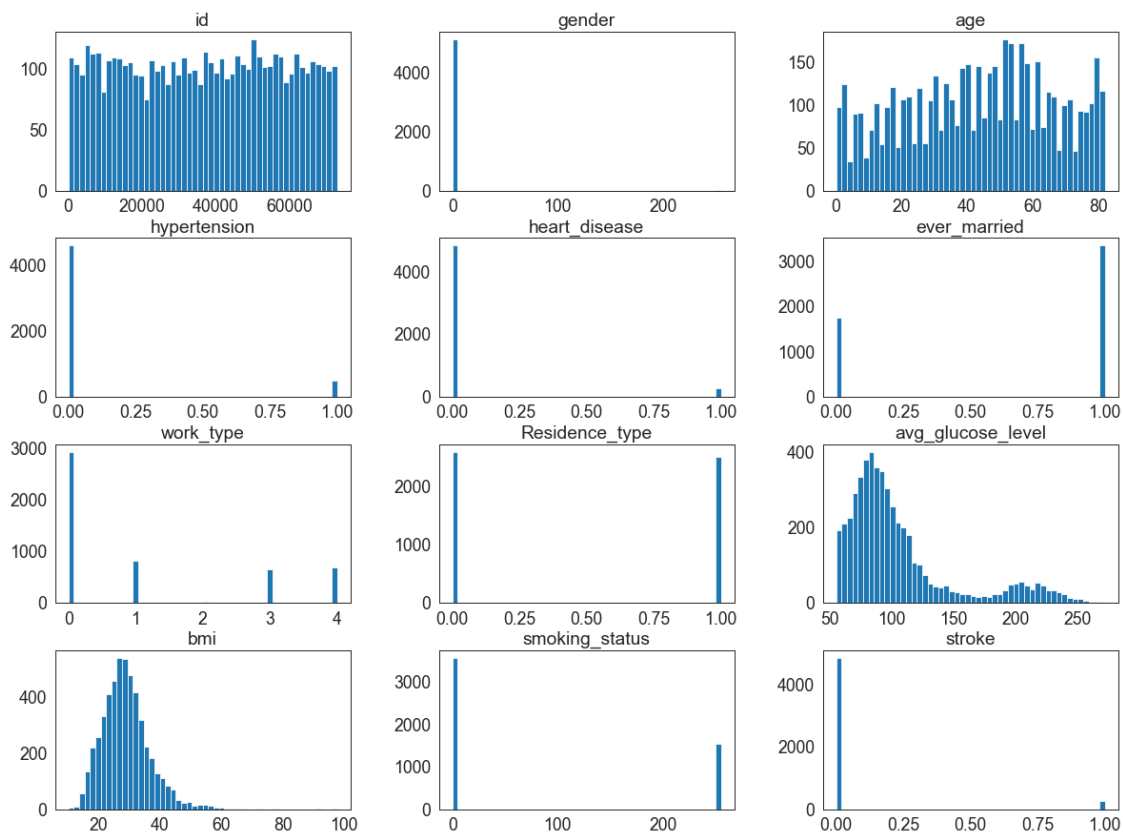
	stroke
0	1
1	1
2	1
3	1
4	1
...	...
5105	0
5106	0
5107	0
5108	0
5109	0

[5110 rows x 12 columns]

## 2.2 plot the distribution of every predictors with histogram

```
[9]: df_stroke_int.hist(bins = 50,grid = False, figsize = (20,15))
```


```
[9]: array([[<AxesSubplot:title={'center':'id'}>,  
          <AxesSubplot:title={'center':'gender'}>,  
          <AxesSubplot:title={'center':'age'}>],  
        [<AxesSubplot:title={'center':'hypertension'}>,  
          <AxesSubplot:title={'center':'heart_disease'}>,  
          <AxesSubplot:title={'center':'ever_married'}>],  
        [<AxesSubplot:title={'center':'work_type'}>,  
          <AxesSubplot:title={'center':'Residence_type'}>,  
          <AxesSubplot:title={'center':'avg_glucose_level'}>],  
        [<AxesSubplot:title={'center':'bmi'}>,  
          <AxesSubplot:title={'center':'smoking_status'}>,  
          <AxesSubplot:title={'center':'stroke'}>]], dtype=object)
```



## 2.3 plot the heat map

```
[10]: f, ax = plt.subplots(figsize = (20,20))
sns.heatmap(df_stroke_int.corr(), annot=True, fmt=".3f", linewidths=0.8,
            ↪ax=ax,cmap="YlGnBu")
```

[10]: <AxesSubplot:>



According to the upper picture, we found that the propability of stroke has little related to work\_type, Residence\_type and smoking\_status.

Let's go further.

```

[11]: df_isstroke = df_stroke[df_stroke['stroke'] == 1]
      df_notstroke = df_stroke[df_stroke['stroke'] != 1]

[12]: fig = plt.figure(figsize=(22,15))
      gs = fig.add_gridspec(3, 3)
      gs.update(wspace=0.35, hspace=0.27)
      ax0 = fig.add_subplot(gs[0, 0])
      ax1 = fig.add_subplot(gs[0, 2])
      ax2 = fig.add_subplot(gs[1, 1])
      ax3 = fig.add_subplot(gs[2, 0])
      ax4 = fig.add_subplot(gs[2, 2])
      background_color = "#f6f6f6"
      fig.patch.set_facecolor(background_color) # figure background color

      #age
      ax0.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
      positive = pd.DataFrame(df_isstroke["age"])
      negative = pd.DataFrame(df_notstroke["age"])
      sns.kdeplot(positive["age"], ax=ax0, color="#0f4c81", shade=True,
        ↳ ec='black', label="positive")
      sns.kdeplot(negative["age"], ax=ax0, color="#9bb7d4", shade=True,
        ↳ ec='black', label="negative")
      ax0.yaxis.set_major_locator(mtick.MultipleLocator(2))
      ax0.set_ylabel('')
      ax0.set_xlabel('')
      ax0.text(-20, 0.0465, 'Age', fontsize=14, fontweight='bold',
        ↳ fontfamily='serif', color="#323232")

      #gender
      positive = pd.DataFrame(df_isstroke["gender"].value_counts())
      positive["Percentage"] = positive["gender"].apply(lambda x: x/
        ↳ sum(positive["gender"])*100)
      negative = pd.DataFrame(df_notstroke["gender"].value_counts())
      negative["Percentage"] = negative["gender"].apply(lambda x: x/
        ↳ sum(negative["gender"])*100)

      x = np.arange(len(positive))
      ax1.text(-0.4, 68.5, 'Gender', fontsize=14, fontweight='bold',
        ↳ fontfamily='serif', color="#323232")
      ax1.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
      ax1.bar(x, height=positive["Percentage"], zorder=3, color="#0f4c81", width=0.4)

      ####problem

      ax1.bar(x+0.4, height=negative.loc[['Female', 'Male'],]["Percentage"], zorder=3,
        ↳ color="#9bb7d4", width=0.4)
      ax1.set_xticks(x + 0.4 / 2)
      ax1.set_xticklabels(['Male', 'Female'])

```



```

ax1.yaxis.set_major_formatter(mtick.PercentFormatter())
ax1.yaxis.set_major_locator(mtick.MultipleLocator(10))
for i,j in zip([0, 1], positive["Percentage"]):
    ax1.annotate(f'{j:0.0f}%',xy=(i, j/2), color='#f6f6f6',
    ↪horizontalalignment='center', verticalalignment='center')
for i,j in zip([0, 1], negative["Percentage"]):
    ax1.annotate(f'{j:0.0f}%',xy=(i+0.4, j/2), color='#f6f6f6',
    ↪horizontalalignment='center', verticalalignment='center')

#Avg. Glucose level
ax2.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
positive = pd.DataFrame(df_isstroke["avg_glucose_level"])
negative = pd.DataFrame(df_notstroke["avg_glucose_level"])
sns.kdeplot(positive["avg_glucose_level"], ax=ax2,color="#0f4c81",ec='black',
    ↪shade=True, label="positive")
sns.kdeplot(negative["avg_glucose_level"], ax=ax2, color="#9bb7d4",
    ↪ec='black',shade=True, label="negative")
ax2.text(-55, 0.01855, 'Avg. Glucose Level',
    fontsize=14, fontweight='bold', fontfamily='serif', color="#323232")
ax2.yaxis.set_major_locator(mtick.MultipleLocator(2))
ax2.set_ylabel('')
ax2.set_xlabel('')

#BMI
ax3.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
positive = pd.DataFrame(df_isstroke["bmi"])
negative = pd.DataFrame(df_notstroke["bmi"])
sns.kdeplot(positive["bmi"], ax=ax3,color="#0f4c81", ec='black',shade=True,
    ↪label="positive")
sns.kdeplot(negative["bmi"], ax=ax3, color="#9bb7d4",ec='black', shade=True,
    ↪label="negative")
ax3.text(-0.06, 0.09, 'BMI',
    fontsize=14, fontweight='bold', fontfamily='serif', color="#323232")
ax3.yaxis.set_major_locator(mtick.MultipleLocator(2))
ax3.set_ylabel('')
ax3.set_xlabel('')

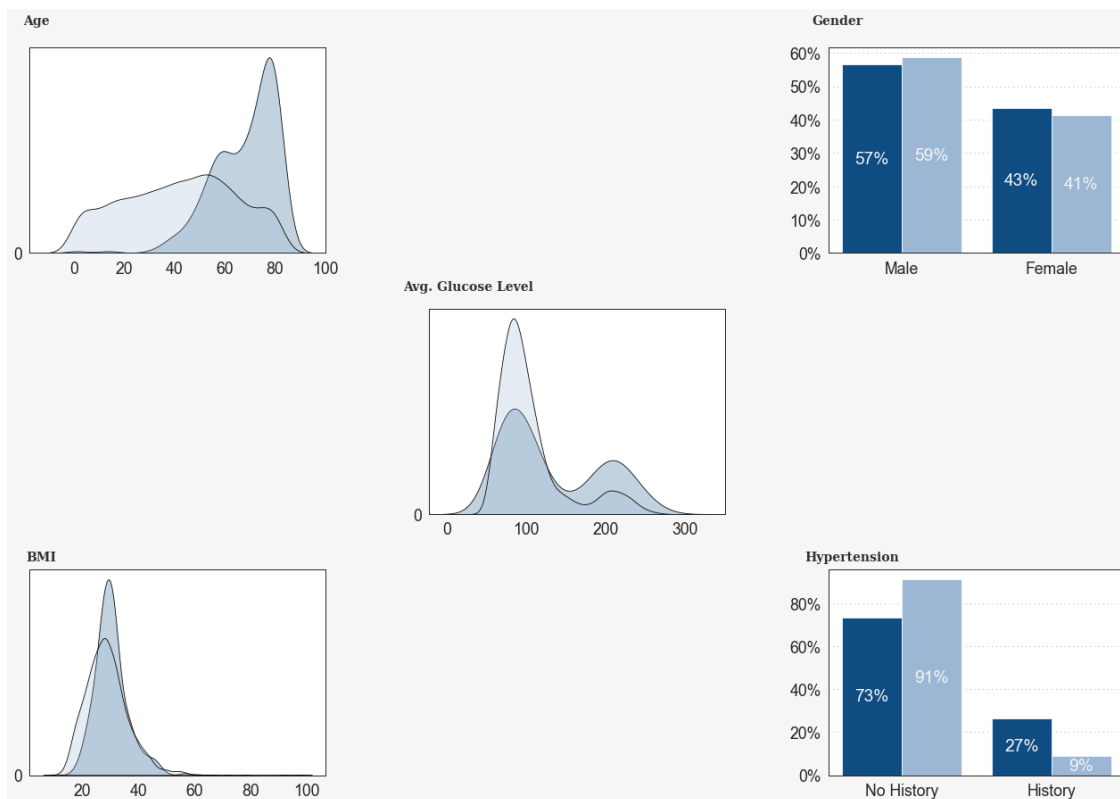
#Hypertension
positive = pd.DataFrame(df_isstroke["hypertension"].value_counts())
positive["Percentage"] = positive["hypertension"].apply(lambda x: x/
    ↪sum(positive["hypertension"])*100)
negative = pd.DataFrame(df_notstroke["hypertension"].value_counts())
negative["Percentage"] = negative["hypertension"].apply(lambda x: x/
    ↪sum(negative["hypertension"])*100)

```

```

x = np.arange(len(positive))
ax4.text(-0.45, 100, 'Hypertension', fontsize=14, fontweight='bold',
        fontfamily='serif', color="#323232")
ax4.grid(color='gray', linestyle=':', axis='y', zorder=0, dashes=(1,5))
ax4.bar(x, height=positive["Percentage"], zorder=3, color="#0f4c81", width=0.4)
ax4.bar(x+0.4, height=negative["Percentage"], zorder=3, color="#9bb7d4",
        width=0.4)
ax4.set_xticks(x + 0.4 / 2)
ax4.set_xticklabels(['No History', 'History'])
ax4.yaxis.set_major_formatter(mtick.PercentFormatter())
ax4.yaxis.set_major_locator(mtick.MultipleLocator(20))
for i,j in zip([0, 1], positive["Percentage"]):
    ax4.annotate(f'{j:0.0f}%', xy=(i, j/2), color='#f6f6f6',
                horizontalalignment='center', verticalalignment='center')
for i,j in zip([0, 1], negative["Percentage"]):
    ax4.annotate(f'{j:0.0f}%', xy=(i+0.4, j/2), color='#f6f6f6',
                horizontalalignment='center', verticalalignment='center')

```



So we build model with upper predictors

### 3 Step three: Before building models, we affirm the training and testing datas first.

#### 3.1 We first look at the number of the data who has stroke and who has't

```
[13]: nbisstroke = len(df_isstroke)
      nbnotstroke = len(df_notstroke)
      print('number of people who has stroked',nbisstroke)
      print("number of people who hasn't stroked",nbnotstroke)
```

number of people who has stroked 249  
number of people who hasn't stroked 4861

```
[14]: fig = px.pie(df_stroke,names='stroke')
      fig.update_layout(title='dataset <b>stroke or not</b> Propotion')
      fig.show()
```

dataset **stroke or not** Propotion



### We find that there are too many data of who hasn't attacked by stroke, so we should balance these numbers of datas.

We now solve over-sampling problem.

#### 3.2 We use SMOTE(Synthetic Minority Over-sampling Technique) to generate more datas.

The traditional over-sampling methode, just duplicating these minority sample and making no new infomation to the dataset, which will easily causes over-fitting problem. SMOTE works by selecting examples that are close in the feature space, drawing a line between the examples in the feature space and drawing a new sample at a point along that line. Aliking the equation:  $x_{new} = x + rand(0,1) * (\hat{x} - x)$

```
[15]: oversample = SMOTE()
      df_precise =_
      ↪df_stroke[['gender','age','hypertension','avg_glucose_level','bmi','stroke']]
```

```

df_precise.gender = df_precise.gender.replace({'Male':0,'Female':1,'Other':-1}).
↳astype(np.uint8)
df_test = df_precise.sample(int(len(df_stroke)*0.2),random_state = 30)
df_train = df_precise.drop(index = df_test.index)

x_test, y_test =↳
↳df_test[['gender','age','hypertension','avg_glucose_level','bmi']],↳
↳df_test['stroke']
x_train, y_train =↳
↳df_train[['gender','age','hypertension','avg_glucose_level','bmi']],↳
↳df_train['stroke']

x_test, y_test = oversample.fit_resample(x_test,y_test)
adjust_test = x_test.assign(stroke = y_test)
x_train, y_train = oversample.fit_resample(x_train,y_train)
adjust_train = x_train.assign(stroke = y_train)

```

```

[16]: print(len(adjust_test[adjust_test.stroke == 1]),len(adjust_test[adjust_test.
↳stroke == 0]))
print(len(adjust_train[adjust_train.stroke == 1]),len(adjust_train[adjust_train.
↳stroke == 0]))

```


986 986  
3875 3875

```

[17]: fig = px.pie(adjust_train,names='stroke')
fig.update_layout(title='train data <b>stroke or not</b> Propotion after↳
↳upsambling')
fig.show()

```

train data **stroke or not** Propotion after upsampling




```

[18]: fig = px.pie(adjust_test,names='stroke')
fig.update_layout(title='test data <b>stroke or not</b> Propotion after↳
↳upsambling')

```

```
fig.show()
```

test data **stroke or not** Proportion after upsampling



## 4 Step four: Build models and analyse the result

I will try following models:

Random Forest

Logistic Regression

Decision Tree

MLP Neural Network with L-BFGS

MLP Neural Network with SGD

Support Vector Machine

```
[19]: # model
RF_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('RF',RandomForestClassifier(random_state = 30))
])
LogR_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('LogR',LogisticRegression(random_state = 30))
])
DT_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('DT',DecisionTreeClassifier(random_state = 30))
])
MLP_lbfgs_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('MLP',MLPClassifier(solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(5, 2),
    random_state = 30))
```

```

])
MLP_SGD_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('MLP',MLPClassifier(hidden_layer_sizes=(5,2), random_state=30,
    ↪max_iter=40, warm_start=True))
])
SVM_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('SVM',SVC(random_state = 30))
])

```

```

[20]: # cross validation
RF_CV = cross_val_score(RF_pipeline, x_train, y_train, cv = 10, scoring = 'f1')
LogR_CV = cross_val_score(LogR_pipeline, x_train, y_train, cv = 10, scoring =
    ↪'f1')
DT_CV = cross_val_score(DT_pipeline, x_train, y_train, cv = 10, scoring = 'f1')
MLP_lbfgs_CV = cross_val_score(MLP_lbfgs_pipeline, x_train, y_train, cv = 10,
    ↪scoring = 'f1')
MLP_SGD_CV = cross_val_score(MLP_SGD_pipeline, x_train, y_train, cv = 10,
    ↪scoring = 'f1')
SVM_CV = cross_val_score(SVM_pipeline, x_train, y_train, cv = 10, scoring =
    ↪'f1')

```

```

[21]: print('The correction rate of cross validation of Random Forest: ', RF_CV.
    ↪mean())
print('The correction rate of cross validation of Logistic Regression: ',
    ↪LogR_CV.mean())
print('The correction rate of cross validation of Decision Tree: ', DT_CV.
    ↪mean())
print('The correction rate of cross validation of MultiLayer Perception with
    ↪L-BFGS algorithm: ', MLP_lbfgs_CV.mean())
print('The correction rate of cross validation of MultiLayer Perception with
    ↪Stochastic Gradient Descent algorithm: ', MLP_SGD_CV.mean())
print('The correction rate of cross validation of Support Vector Machine : ',
    ↪SVM_CV.mean())

```

The correction rate of cross validation of Random Forest: 0.9254277042396863

The correction rate of cross validation of Logistic Regression:

0.8125404112385366

The correction rate of cross validation of Decision Tree: 0.9000791816875742

The correction rate of cross validation of MultiLayer Perception with L-BFGS  
algorithm: 0.8193992873591979

The correction rate of cross validation of MultiLayer Perception with Stochastic  
Gradient Descent algorithm: 0.81339693090295

The correction rate of cross validation of Support Vector Machine :

0.8224945320864562

From what we have seen above, we can conclusion easily that the model Random Forest perform

best.

But one who is ill'd wouldn't like to get a wrong answer, to make it more reasonable, we next calculate the Recall Rate

```
[22]: RF_pipeline.fit(x_train,y_train)
      LogR_pipeline.fit(x_train,y_train)
      DT_pipeline.fit(x_train,y_train)
      MLP_lbfgs_pipeline.fit(x_train,y_train)
      MLP_SGD_pipeline.fit(x_train,y_train)
      SVM_pipeline.fit(x_train,y_train)
```

```
[22]: Pipeline(steps=[('scale', StandardScaler()), ('SVM', SVC(random_state=30))])
```

```
[23]: # predict and calculate confusion-matrix and f1-score
      RF_pred = RF_pipeline.predict(x_test)
      LogR_pred = LogR_pipeline.predict(x_test)
      DT_pred = DT_pipeline.predict(x_test)
      MLP_lbfgs_pred = MLP_lbfgs_pipeline.predict(x_test)
      MLP_SGD_pred = MLP_SGD_pipeline.predict(x_test)
      SVM_pred = SVM_pipeline.predict(x_test)

      RF_conf = confusion_matrix(y_test, RF_pred)
      LogR_conf = confusion_matrix(y_test, LogR_pred)
      DT_conf = confusion_matrix(y_test, DT_pred)
      MLP_lbfgs_conf = confusion_matrix(y_test, MLP_lbfgs_pred)
      MLP_SGD_conf = confusion_matrix(y_test, MLP_SGD_pred)
      SVM_conf = confusion_matrix(y_test, SVM_pred)

      RF_f1 = f1_score(y_test, RF_pred)
      LogR_f1 = f1_score(y_test, LogR_pred)
      DT_f1 = f1_score(y_test, DT_pred)
      MLP_lbfgs_f1 = f1_score(y_test, MLP_lbfgs_pred)
      MLP_SGD_f1 = f1_score(y_test, MLP_SGD_pred)
      SVM_f1 = f1_score(y_test, SVM_pred)

      RF_acc = accuracy_score(y_test, RF_pred)
      LogR_acc = accuracy_score(y_test, LogR_pred)
      DT_acc = accuracy_score(y_test, DT_pred)
      MLP_lbfgs_acc = accuracy_score(y_test, MLP_lbfgs_pred)
      MLP_SGD_acc = accuracy_score(y_test, MLP_SGD_pred)
      SVM_acc = accuracy_score(y_test, SVM_pred)

      RF_rec = recall_score(y_test, RF_pred)
      LogR_rec = recall_score(y_test, LogR_pred)
      DT_rec = recall_score(y_test, DT_pred)
      MLP_lbfgs_rec = recall_score(y_test, MLP_lbfgs_pred)
      MLP_SGD_rec = recall_score(y_test, MLP_SGD_pred)
```

```

SVM_rec = recall_score(y_test, SVM_pred)


RF_pre = precision_score(y_test, RF_pred)
LogR_pre = precision_score(y_test, LogR_pred)
DT_pre = precision_score(y_test, DT_pred)
MLP_lbfgs_pre = precision_score(y_test, MLP_lbfgs_pred)
MLP_SGD_pre = precision_score(y_test, MLP_SGD_pred)
SVM_pre = precision_score(y_test, SVM_pred)

```

```

[24]: # f1-score
fig = pgo.Figure()
fig.add_trace(pgo.Bar(x = ['Random Forest', 'Logistic Regression', 'Decision_
↳Tree', 'MLP Neural Network with L-BFGS', 'MLP Neural Network with_
↳SGD', 'Support Vector Machine'], y =_
↳[RF_f1, LogR_f1, DT_f1, MLP_lbfgs_f1, MLP_SGD_f1, SVM_f1]))
fig.update_layout(title = 'f1-score of every model', xaxis_title = 'Model_
↳Name', yaxis_title = 'f1-score')

```




```

[25]: # Random Forest
fig = ff.create_annotated_heatmap(RF_conf, x=['Not Stroke', 'Is Stroke'],_
↳y=['Predicted Not Stroke', 'Predicted Is Stroke'],_
↳colorscale='blugrn', xgap=3, ygap=3)
fig['data'][0]['showscale'] = True
fig.update_layout(title='Confusion Matrix with Prediction of Random Forest')
fig.show()
print(classification_report(y_test, RF_pred))
print('Accuracy Score: ', RF_acc)

```





	precision	recall	f1-score	support
0	0.65	0.91	0.76	986
1	0.85	0.52	0.64	986
accuracy			0.71	1972
macro avg	0.75	0.71	0.70	1972
weighted avg	0.75	0.71	0.70	1972

Accuracy Score: 0.7119675456389453


```
[26]: # Logistic Regression
fig = ff.create_annotated_heatmap(LogR_conf, x=['Not Stroke','Is Stroke'],
    ↳y=['Predicted Not Stroke','Predicted Is Stroke'],
    ↳colorscale='blugrn',xgap=3,ygap=3)
fig['data'][0]['showscale'] = True
fig.update_layout(title='Confusion Matrix with Prediction of Logistic
    ↳Regression')
fig.show()
print(classification_report(y_test,LogR_pred))
print('Accuracy Score: ', LogR_acc)
```



	precision	recall	f1-score	support
0	0.75	0.73	0.74	986
1	0.74	0.76	0.75	986
accuracy			0.75	1972
macro avg	0.75	0.75	0.75	1972
weighted avg	0.75	0.75	0.75	1972

Accuracy Score: 0.7454361054766734


```
[27]: # Decision Tree
fig = ff.create_annotated_heatmap(DT_conf, x=['Not Stroke','Is Stroke'],
    ↳y=['Predicted Not Stroke','Predicted Is Stroke'],
    ↳colorscale='blugrn',xgap=3,ygap=3)
fig['data'][0]['showscale'] = True
fig.update_layout(title='Confusion Matrix with Prediction of Decision Tree')
fig.show()
print(classification_report(y_test,DT_pred))
print('Accuracy Score: ', DT_acc)
```



	precision	recall	f1-score	support
0	0.61	0.89	0.73	986
1	0.80	0.44	0.57	986
accuracy			0.66	1972
macro avg	0.71	0.66	0.65	1972
weighted avg	0.71	0.66	0.65	1972

Accuracy Score: 0.6648073022312373


```
[28]: # MLP Neural Network with L-BFGS
fig = ff.create_annotated_heatmap(MLP_lbfgs_conf, x=['Not Stroke','Is Stroke'],
    y=['Predicted Not Stroke','Predicted Is Stroke'],
    colorscale='blugrn',xgap=3,ygap=3)
fig['data'][0]['showscale'] = True
fig.update_layout(title='Confusion Matrix with Prediction of MLP Neural Network
    with L-BFGS')
fig.show()
print(classification_report(y_test,MLP_lbfgs_pred))
print('Accuracy Score: ', MLP_lbfgs_acc)
```



	precision	recall	f1-score	support
0	0.72	0.76	0.74	986
1	0.75	0.70	0.73	986
accuracy			0.73	1972
macro avg	0.73	0.73	0.73	1972
weighted avg	0.73	0.73	0.73	1972

Accuracy Score: 0.7327586206896551


```
[29]: # MLP Neural Network with SGD
fig = ff.create_annotated_heatmap(MLP_SGD_conf, x=['Not Stroke','Is Stroke'],
    y=['Predicted Not Stroke','Predicted Is Stroke'],
    colorscale='blugrn',xgap=3,ygap=3)
fig['data'][0]['showscale'] = True
fig.update_layout(title='Confusion Matrix with Prediction of MLP Neural Network
    with SGD')
fig.show()
print(classification_report(y_test,MLP_SGD_pred))
print('Accuracy Score: ', MLP_SGD_acc)
```



	precision	recall	f1-score	support
0	0.82	0.71	0.76	986
1	0.74	0.84	0.79	986
accuracy			0.77	1972
macro avg	0.78	0.77	0.77	1972
weighted avg	0.78	0.77	0.77	1972

Accuracy Score: 0.7748478701825557

```
[30]: # Support Vector Machine
fig = ff.create_annotated_heatmap(SVM_conf, x=['Not Stroke','Is Stroke'],
    ↳y=['Predicted Not Stroke','Predicted Is Stroke'],
    ↳colorscale='blugrn',xgap=3,ygap=3)
fig['data'][0]['showscale'] = True
fig.update_layout(title='Confusion Matrix with Prediction of Support Vector
    ↳Machine')
fig.show()
print(classification_report(y_test,SVM_pred))
print('Accuracy Score: ', SVM_acc)
```



	precision	recall	f1-score	support
0	0.78	0.72	0.75	986
1	0.74	0.80	0.77	986
accuracy			0.76	1972
macro avg	0.76	0.76	0.76	1972
weighted avg	0.76	0.76	0.76	1972

Accuracy Score: 0.7586206896551724

## 5 Step five: Model Selection

```
[31]: # make dataframes to plot
RF_df = pd.DataFrame(data = [RF_f1,RF_acc,RF_rec,RF_pre], columns = ['Random_
    ↳Forest'], index = ['f1','accuracy','recall','precision'])
LogR_df = pd.DataFrame(data = [LogR_f1,LogR_acc,LogR_rec,LogR_pre], columns =
    ↳['Logistic Regression'], index = ['f1','accuracy','recall','precision'])
DT_df = pd.DataFrame(data = [DT_f1,DT_acc,DT_rec,DT_pre], columns = ['Decision_
    ↳Tree'], index = ['f1','accuracy','recall','precision'])
MLP_lbfgs_df = pd.DataFrame(data =
    ↳[MLP_lbfgs_f1,MLP_lbfgs_acc,MLP_lbfgs_rec,MLP_lbfgs_pre], columns = ['MLP_
    ↳Neural Network with L-BFGS'], index = ['f1','accuracy','recall','precision'])
MLP_SGD_df = pd.DataFrame(data =
    ↳[MLP_SGD_f1,MLP_SGD_acc,MLP_SGD_rec,MLP_SGD_pre], columns = ['MLP Neural_
    ↳Network with SGD'], index = ['f1','accuracy','recall','precision'])
SVM_df = pd.DataFrame(data = [SVM_f1,SVM_acc,SVM_rec,SVM_pre], columns =
    ↳['Support Vector Machine'], index = ['f1','accuracy','recall','precision'])

[32]: df_analysis = round(pd.
    ↳concat([RF_df,LogR_df,DT_df,MLP_lbfgs_df,MLP_SGD_df,SVM_df], axis=1) , 6)
colors = ["lightgray","lightgray","#0f4c81"]
colormap = matplotlib.colors.LinearSegmentedColormap.from_list("", colors)
background_color = "#fbfbfb"
fig = plt.figure(figsize=(18,16))
gs = fig.add_gridspec(3, 5)
gs.update(wspace=0.1, hspace=0.5)
ax0 = fig.add_subplot(gs[0, :])
sns.heatmap(df_analysis.T, cmap=colormap,annot=True,fmt=".1%",vmin=0,vmax=0.95,
    ↳linewidths=2.5,cbar=False,ax=ax0,annot_kws={"fontsize":15})
fig.patch.set_facecolor(background_color) # figure background color
ax0.set_facecolor(background_color)
ax0.text(1,-0.5,'Model_
    ↳Comparison',fontsize=30,fontweight='bold',fontfamily='Arial Black')
```

```
ax0.tick_params(axis=u'both', which=u'both',length=0)

plt.show()
```

Model Comparison				
Random Forest	64.3%	71.2%	51.8%	84.6%
Logistic Regression	74.9%	74.5%	76.0%	73.9%
Decision Tree	56.8%	66.5%	44.0%	79.9%
MLP Neural Network with L-BFGS	72.5%	73.3%	70.5%	74.7%
MLP Neural Network with SGD	78.9%	77.5%	84.0%	74.3%
Support Vector Machine	76.8%	75.9%	79.9%	73.9%
	f1	accuracy	recall	precision

5.0.1 So obviously, we find that the model of MLP Neural Network with SGD algorithm has the highest recall rate of class\_1 (is\_stroke) which means that it will misclassify least cases of patient who has stroke.

## 6 Step six: Find the best parameters

### 6.1 Find the iterater times of largest recall rate

```
[34]: def MLP_SGD_opt():
    MLP_SGD_rec_max = -1
    MLP_SGD_rec_max_index = -1
    for times in np.arange(1,50):
        MLP_SGD_pipeline = Pipeline(steps = [
            ('scale',StandardScaler()),
            ('MLP',MLPClassifier(hidden_layer_sizes=(5,2), random_state=30,
→max_iter=times, warm_start=True))
        ])
        MLP_SGD_pipeline.fit(x_train,y_train)
        MLP_SGD_pred = MLP_SGD_pipeline.predict(x_test)
        MLP_SGD_rec = recall_score(y_test, MLP_SGD_pred)
    #     print(MLP_SGD_rec, ' ', times)
        if(MLP_SGD_rec_max <= MLP_SGD_rec):
            MLP_SGD_rec_max = MLP_SGD_rec
            MLP_SGD_rec_max_index = times

        else:
            break
    return [MLP_SGD_rec_max,MLP_SGD_rec_max_index]
[m,i] = MLP_SGD_opt()
print("recall rate: ",m,"index: ",i)
```

recall rate: 0.8539553752535497 index: 24

## 6.2 Replot the accuracy rate of this parameter

```
[35]: MLP_SGD_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('MLP',MLPClassifier(hidden_layer_sizes=(5,2), random_state=30,
    ↪max_iter=23, warm_start=True)) #0.981
])

MLP_SGD_pipeline.fit(x_train,y_train)
MLP_SGD_pred = MLP_SGD_pipeline.predict(x_test)
MLP_SGD_conf = confusion_matrix(y_test, MLP_SGD_pred)
MLP_SGD_f1 = f1_score(y_test, MLP_SGD_pred)
MLP_SGD_acc = accuracy_score(y_test, MLP_SGD_pred)
MLP_SGD_rec = recall_score(y_test, MLP_SGD_pred)
MLP_SGD_pre = precision_score(y_test, MLP_SGD_pred)

print("recall rate:",MLP_SGD_rec)

# MLP Neural Network with SGD
fig = ff.create_annotated_heatmap(MLP_SGD_conf, x=['Not Stroke','Is Stroke'],
    ↪y=['Predicted Not Stroke','Predicted Is Stroke'],
    ↪colorscale='blugrn',xgap=3,ygap=3)
fig['data'][0]['showscale'] = True
fig.update_layout(title='Confusion Matrix with Prediction of MLP Neural Network
    ↪with SGD')
fig.show()
print(classification_report(y_test,MLP_SGD_pred))
print('Accuracy Score: ', MLP_SGD_acc)
```

recall rate: 0.8529411764705882



	precision	recall	f1-score	support
0	0.83	0.71	0.76	986

	1	0.74	0.85	0.79	986
accuracy				0.78	1972
macro avg		0.79	0.78	0.78	1972
weighted avg		0.79	0.78	0.78	1972

Accuracy Score: 0.7794117647058824

### 6.3 Conclusion

Due to the randomness of the method stochastic gradient descend and we can do early stopping when we training the data with MultiLayer Network, we can control when to stop by the will of my own. So in this case, I can gain a higher recall rate by sacrificing the accuracy rate. But if we want the best performance of accuracy, (balanced) random forest is still the best choice.

## 7 Last step: Train the model with all data

```
[36]: import pandas as pd
import numpy as np
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor
from imblearn.over_sampling import SMOTE
from sklearn.neural_network import MLPClassifier
import warnings
warnings.filterwarnings('ignore')
```

```
[37]: df_stroke_all = pd.read_csv("../data/healthcare-dataset-stroke-data.csv")
bmi_pipe = Pipeline( steps = [
    ('scaling',StandardScaler()),
    ('lr',DecisionTreeRegressor(random_state = 30))
])
cp = df_stroke_all[['age','gender','bmi']].copy()
cp.gender = cp.gender.replace({'Male':0,'Female':1,'Other':-1}).astype(np.uint8)
miss_bmi = cp[cp.bmi.isnull()]
cp = cp[~cp.bmi.isnull()]
bmi = cp.pop('bmi')
bmi_pipe.fit(cp,bmi)
predict_bmi = pd.Series(bmi_pipe.predict(miss_bmi[['age','gender']] ),index =
    ↳miss_bmi.index)
df_stroke_all.loc[miss_bmi.index,'bmi'] = predict_bmi

oversample = SMOTE()
df_precise_all =
    ↳df_stroke_all[['gender','age','hypertension','avg_glucose_level','bmi','stroke']]
df_precise_all.gender = df_precise_all.gender.replace({'Male':0,'Female':
    ↳1,'Other':-1}).astype(np.uint8)
```



```

x_all, y_all =
↳df_precise_all[['gender','age','hypertension','avg_glucose_level','bmi']],
↳df_precise_all['stroke']
x_all, y_all = oversample.fit_resample(x_all,y_all)
adjust_all = x_all.assign(stroke = y_all)
MLP_SGD_pipeline = Pipeline(steps = [
    ('scale',StandardScaler()),
    ('MLP',MLPClassifier(hidden_layer_sizes=(5,2), random_state=30,
↳max_iter=40, warm_start=True))
])
MLP_SGD_pipeline.fit(x_all,y_all)

```

```

[37]: Pipeline(steps=[('scale', StandardScaler()),
    ('MLP',
    MLPClassifier(hidden_layer_sizes=(5, 2), max_iter=40,
    random_state=30, warm_start=True))])

```

## 7.1 Test with several data

```

[38]: MLP_SGD_pipeline.predict(pd.DataFrame([[0,67,0,228,36.6]]))

```

```

[38]: array([1])

```

```

[39]: MLP_SGD_pipeline.predict(pd.DataFrame([[1,79,1,83.75,28.4]]))

```

```

[39]: array([1])

```

## 7.2 Store the optimal model

```

[40]: joblib.dump(MLP_SGD_pipeline,'./data/model_MLP_SGD_fitted.pkl')

```

```

[40]: ['./data/model_MLP_SGD_fitted.pkl']

```

```

[41]: # load the model

# model = joblib.load("./data/model_MLP_SGD_fitted.pkl")

# model.predict(pd.DataFrame([[1,79,1,83.75,28.4]]))

```

```

[ ]:

```